Inference Is All You Need

-

What did Ilya see?

Sam Altman
@sama · Follow

i love summer in the garden

5:29 PM · Aug 7, 2024

12.5K Reply Share

Read 1.4K replies

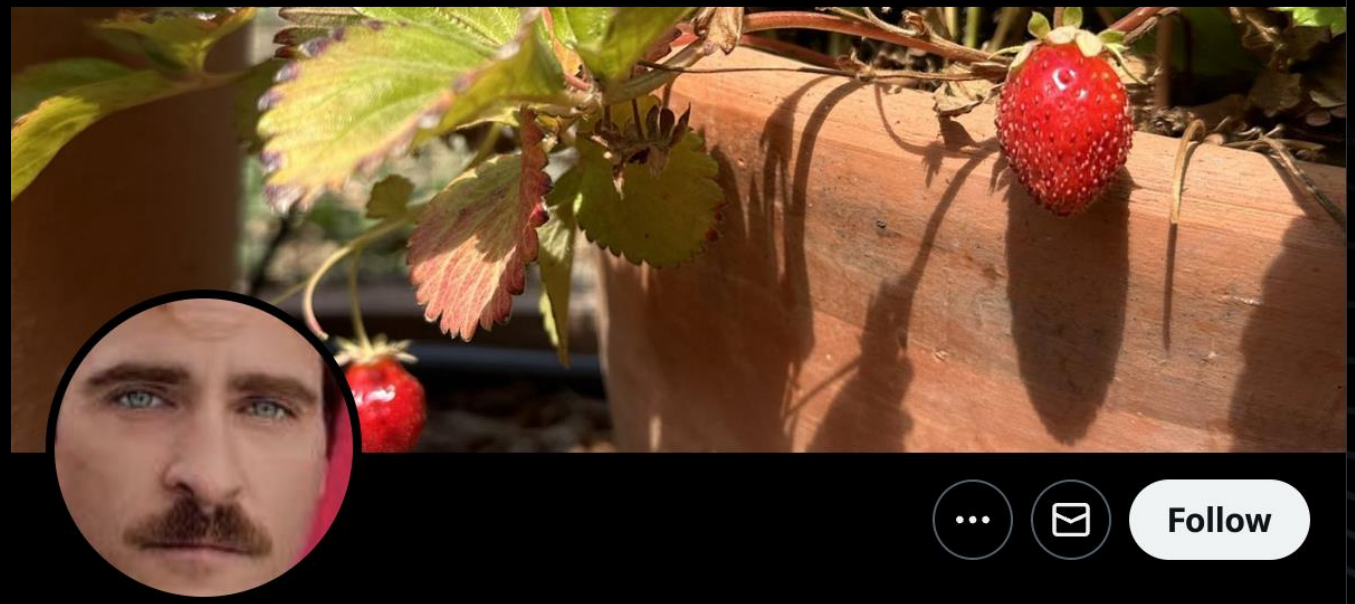

Follow

@iruletheworldmo

purely theatre

Science & Technology Joined August 2022

1,673 Following **30K** Followers

ellamind

Large Language Monkeys: Scaling Inference Compute with Repeated Sampling

Bradley Brown*†‡, Jordan Juravsky*†, Ryan Ehrlich*†, Ronald Clark‡, Quoc V. Le§, Christopher Ré†, and Azalia Mirhoseini†§

†Department of Computer Science, Stanford University
‡University of Oxford
§Google DeepMind

bradley.brown@cs.ox.ac.uk, jbj@stanford.edu, ryanehrlich@cs.stanford.edu, ronald.clark@cs.ox.ac.uk, qvl@google.com, chrismre@stanford.edu, azalia@stanford.edu

July 30, 2024

Google DeepMind 2024-8-7

Scaling LLM Test-Time Compute Optimally can be More Effective than Scaling Model Parameters

Charlie Snell[♦, 1], **Jaehoon Lee**[2], **Kelvin Xu**[♣, 2] **and Aviral Kumar**[♣, 2]

[♣]Equal advising, [1]UC Berkeley, [2]Google DeepMind, [♦]Work done during an internship at Google DeepMind

1. What does ”Scaling Test Time Compute” mean?

2. Why does it matter for AI Engineers?

“Ask the same question a couple of times“

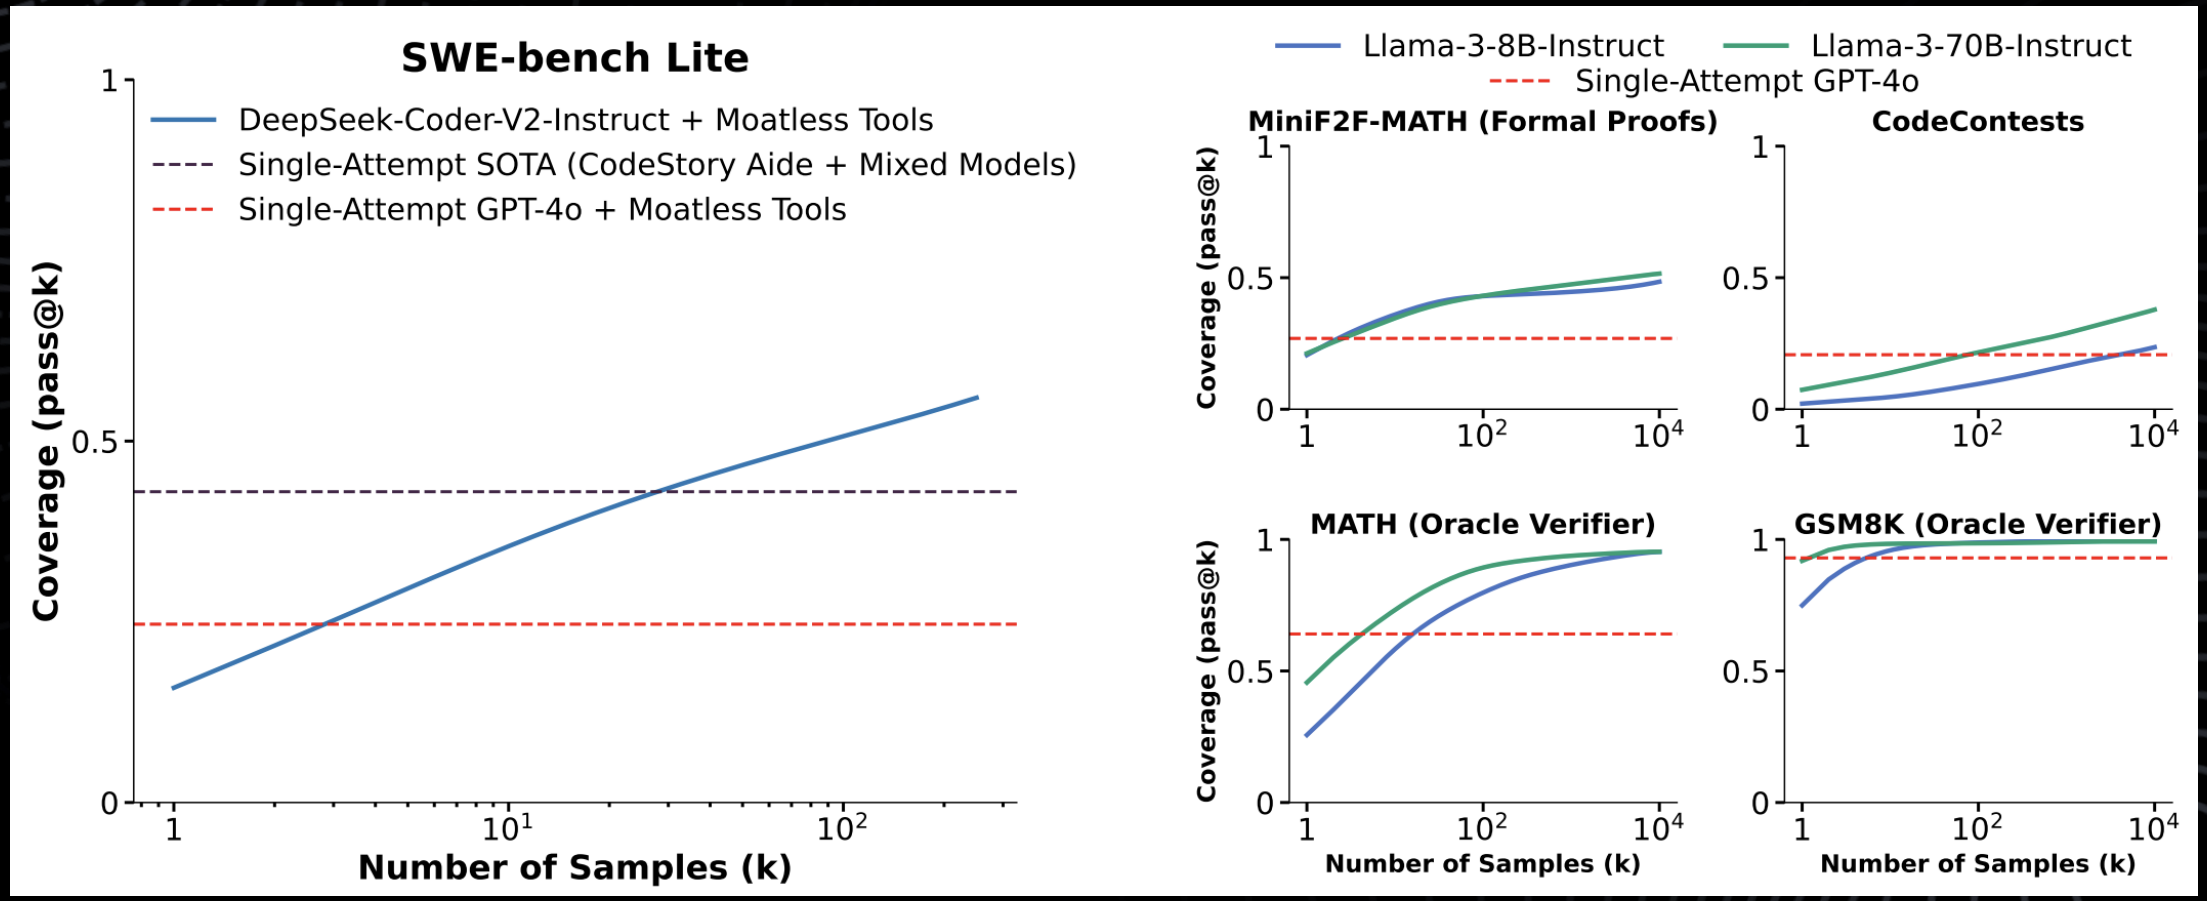

“Ask the same question a couple of times“

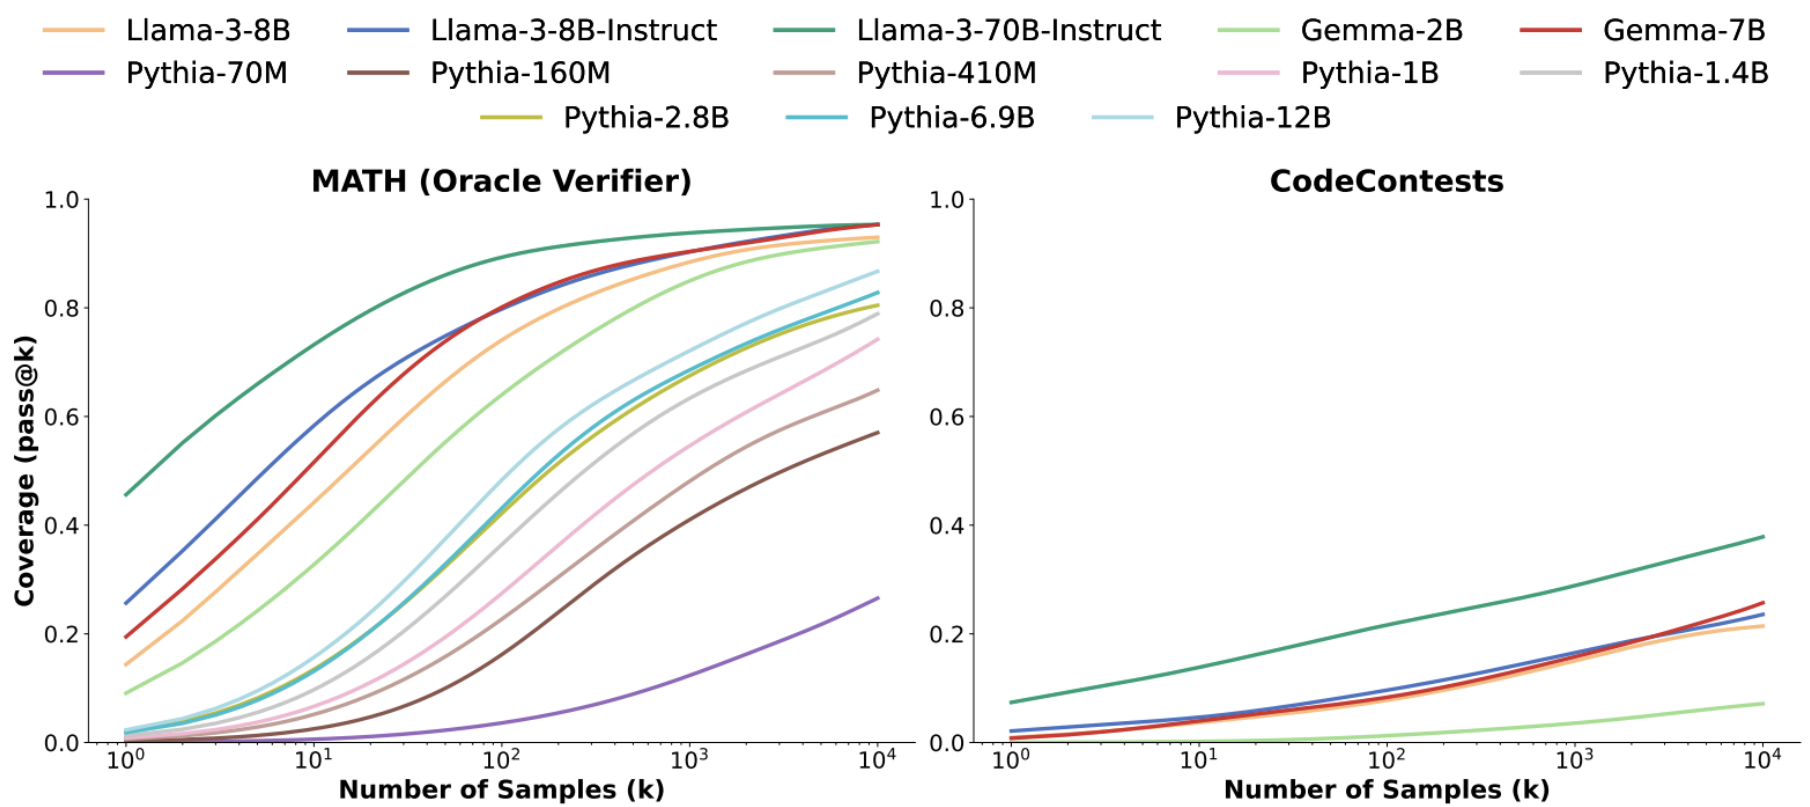

Figure 3: Scaling inference time compute via repeated sampling leads to consistent coverage gains across a variety of model sizes (70M-70B), families (Llama, Gemma and Pythia) and levels of post-training (Base and Instruct models).

Access to “Ground Truth”

Golden Answer
Unit Tests
Formal Verfication
Any strict test

Reward Models!

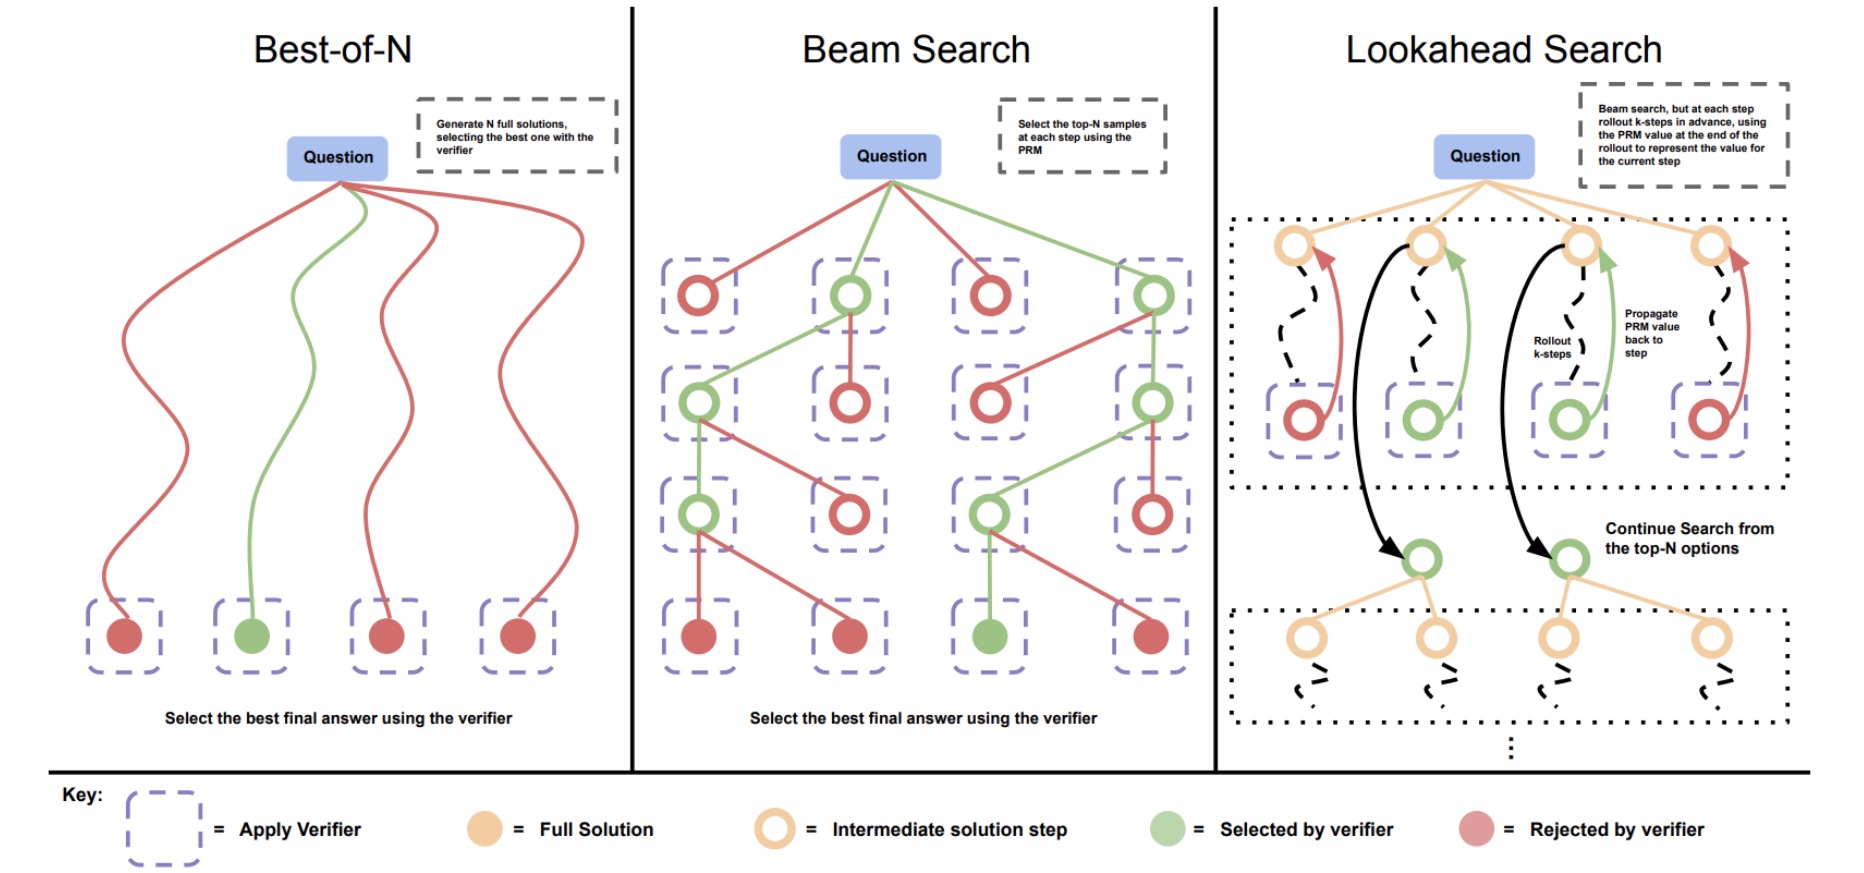

Figure 2 | ***Comparing different PRM search methods.*** **Left:** Best-of-N samples N full answers and then selects the best answer according to the PRM final score. **Center:** Beam search samples N candidates at each step, and selects the top M according to the PRM to continue the search from. **Right:** lookahead-search extends each step in beam-search to utilize a k-step lookahead while assessing which steps to retain and continue the search from. Thus lookahead-search needs more compute.

AGI achieved internally? 🍓

PRM800K

Let's Verify Step by Step

Hunter Lightman* Vineet Kosaraju* Yura Burda* Harri Edwards

Bowen Baker Teddy Lee Jan Leike John Schulman Ilya Sutskever

Karl Cobbe*

OpenAI

Abstract

Why does this matter for an AI engineer?

- **Sample on (idle) low VRAM hardware**
- **Inference is getting more optimized**
  - Groq, Cerebras
  - Quantization
  - Speculative Decoding
- **Collecting data for Instuction Finetuning can be harder than collecting preference data for RM**

rasdani/inference-is-all-you-need

jan@ellamind.com

https://ellamind.com

ellamind GmbH
Konsul-Smidt-Straße 8p
28217 Bremen